

## 657a-as1-ques1

February 9, 2023

```
[124]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
df = pd.read_csv("D:/UWaterloo/Data Knowledge and Modelling/Assignment 1/
↳ abalone.csv", names = ['Sex', 'Length', 'Diameter', 'Height', 'Whole_weight',
                          'Sucked_weight', 'Viscera_weight', 'Shell_weight',
↳ 'Rings'], sep = ',')
warnings.filterwarnings('ignore')
```

```
[125]: df.head()
```

```
[125]:
```

	Sex	Length	Diameter	Height	Whole_weight	Sucked_weight	Viscera_weight	\
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	

	Shell_weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9
3	0.155	10
4	0.055	7

```
[126]: df.isna().any()
```

```
[126]: Sex                False
Length                False
Diameter              False
Height                False
Whole_weight          False
Sucked_weight         False
Viscera_weight        False
Shell_weight          False
Rings                 False
dtype: bool
```

There is no missing data in any of the feature columns. Sex is the only categorical data, with the rest of the features being numerical.

```
[127]: len(df)
```

```
[127]: 4177
```

```
[128]: df.columns
```

```
[128]: Index(['Sex', 'Length', 'Diameter', 'Height', 'Whole_weight', 'Sucked_weight',  
        'Viscera_weight', 'Shell_weight', 'Rings'],  
        dtype='object')
```

```
[129]: df.describe()
```

```
[129]:
```

	Length	Diameter	Height	Whole_weight	Sucked_weight	\
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	
mean	0.523992	0.407881	0.139516	0.828742	0.359367	
std	0.120093	0.099240	0.041827	0.490389	0.221963	
min	0.075000	0.055000	0.000000	0.002000	0.001000	
25%	0.450000	0.350000	0.115000	0.441500	0.186000	
50%	0.545000	0.425000	0.140000	0.799500	0.336000	
75%	0.615000	0.480000	0.165000	1.153000	0.502000	
max	0.815000	0.650000	1.130000	2.825500	1.488000	

	Viscera_weight	Shell_weight	Rings
count	4177.000000	4177.000000	4177.000000
mean	0.180594	0.238831	9.933684
std	0.109614	0.139203	3.224169
min	0.000500	0.001500	1.000000
25%	0.093500	0.130000	8.000000
50%	0.171000	0.234000	9.000000
75%	0.253000	0.329000	11.000000
max	0.760000	1.005000	29.000000

```
[130]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4177 entries, 0 to 4176  
Data columns (total 9 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   Sex              4177 non-null   object  
1   Length           4177 non-null   float64  
2   Diameter         4177 non-null   float64  
3   Height           4177 non-null   float64  
4   Whole_weight     4177 non-null   float64  
5   Sucked_weight    4177 non-null   float64  
6   Viscera_weight   4177 non-null   float64
```

```
7   Shell_weight    4177 non-null   float64
8   Rings           4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

```
[131]: df['Sex'].describe()
```

```
[131]: count      4177
unique         3
top            M
freq         1528
Name: Sex, dtype: object
```

```
[132]: df['Sex'].value_counts()
```

```
[132]: M      1528
I       1342
F       1307
Name: Sex, dtype: int64
```

```
[133]: df.median()
```

```
[133]: Length          0.5450
Diameter          0.4250
Height            0.1400
Whole_weight      0.7995
Sucked_weight     0.3360
Viscera_weight    0.1710
Shell_weight      0.2340
Rings             9.0000
dtype: float64
```

```
[134]: df.var()
```

```
[134]: Length          0.014422
Diameter          0.009849
Height            0.001750
Whole_weight      0.240481
Sucked_weight     0.049268
Viscera_weight    0.012015
Shell_weight      0.019377
Rings             10.395266
dtype: float64
```

```
[135]: df.skew()
```

```
[135]: Length          -0.639873
Diameter          -0.609198
```

```

Height          3.128817
Whole_weight    0.530959
Sucked_weight   0.719098
Viscera_weight  0.591852
Shell_weight    0.620927
Rings           1.114102
dtype: float64

```

```
[136]: df.kurtosis()
```

```

[136]: Length          0.064621
Diameter           -0.045476
Height            76.025509
Whole_weight      -0.023644
Sucked_weight     0.595124
Viscera_weight    0.084012
Shell_weight      0.531926
Rings             2.330687
dtype: float64


```

```


[155]: # Check any relationship between Sex and Rings
plt.scatter(y=df['Sex'], x=df['Rings'])

```

```
[155]: <matplotlib.collections.PathCollection at 0x264ba2e7be0>
```



```
[137]: plt.hist(df['Length'], bins=10)
plt.xlabel('Length')
plt.ylabel('Data')
plt.show()
df['Height'].describe()
df[df['Height'] == 0]
```



```
[137]:
```

	Sex	Length	Diameter	Height	Whole_weight	Sucked_weight	\
1257	I	0.430	0.34	0.0	0.428	0.2065	
3996	I	0.315	0.23	0.0	0.134	0.0575	

	Viscera_weight	Shell_weight	Rings
1257	0.0860	0.1150	8
3996	0.0285	0.3505	6

```
[138]: df[['Whole_weight', 'Sucked_weight', 'Viscera_weight', 'Shell_weight']].
        describe()
```

```
[138]:
```


	Whole_weight	Sucked_weight	Viscera_weight	Shell_weight
count	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.828742	0.359367	0.180594	0.238831
std	0.490389	0.221963	0.109614	0.139203
min	0.002000	0.001000	0.000500	0.001500
25%	0.441500	0.186000	0.093500	0.130000
50%	0.799500	0.336000	0.171000	0.234000
75%	1.153000	0.502000	0.253000	0.329000
max	2.825500	1.488000	0.760000	1.005000

1. The number of male fishes is higher than female and infants.
2. For the column of length, the mean is around 0.523 and median is 0.545 (denoted by the 50% value), which shows the distribution must be left skewed due to the presence of infants. The spread is less as SD is less.
3. The min of height feature is 0.000000, which can be an anomaly as other features have valid values. It can be a missing value that needs to be handled. SD is low so the normal distribution is not spread out a lot. It is close to the mean values.
4. Considering the data of all weights in the dataset, they are highly correlated and so we cannot use each weight against another to predict the age of the fish.

Scatter plot : Length vs Whole-weight

```
[139]: plt.scatter(x=df['Length'], y=df['Whole_weight'])
plt.xlabel('Length')
plt.ylabel('Whole_weight')
```


```
[139]: Text(0, 0.5, 'Whole_weight')
```



Height vs Whole weight Analysis: -> There are outliers when height is around 0.5 and height > 1  
-> There is more or less a constant relationship between whole\_weight and height when height is less than 0.2

```
[140]: plt.scatter(x=df['Height'], y=df['Whole_weight'])  
plt.xlabel('Height')  
plt.ylabel('Whole_weight')
```

```
[140]: Text(0, 0.5, 'Whole_weight')
```




Diameter vs Whole weight Analysis: -> There are not a lot of outliers and hence this feature can be used to derive a relationship.

```
[141]: plt.scatter(x=df['Diameter'], y=df['Whole_weight'])  
plt.xlabel('Diameter')  
plt.ylabel('Whole_weight')
```


```
[141]: Text(0, 0.5, 'Whole_weight')
```





```
[142]: sns.pairplot(df, hue='Rings', height=1.5, aspect=2)
```

```
[142]: <seaborn.axisgrid.PairGrid at 0x264b4f79a90>
```



```
[143]: df['Rings'].value_counts()
```

```
[143]: 9      689
      10     634
      8     568
      11    487
      7     391
      12    267
      6     259
      13    203
      14    126
      5     115
      15    103
      16     67
      17     58
      4     57
      18     42
      19     32
      20     26
      3     15
      21     14
      23      9
      22      6
      27      2
      24      2
      1       1
      26      1
      29      1
      2       1
      25      1
      Name: Rings, dtype: int64
```

The above data shows that the dataset is imbalanced as there are more samples available for Rings between 10 - 15. The number of samples in the current dataset is not very high (=4177) and hence, we can perform oversampling in order to balance it.

## 1 Splitting the dataset

```
[144]: independent = df.iloc[:, 0:8]
      dependent = df.iloc[:, 8]
```

```
[145]: from sklearn.model_selection import train_test_split
      ind_train, ind_test, dep_train, dep_test = train_test_split(independent,
      ↪ dependent, test_size = 0.2, random_state=1)

      # Balance the training dataset
      from imblearn.over_sampling import RandomOverSampler
```

```
os = RandomOverSampler(random_state=1)
ind_train_sampled, dep_train_sampled = os.fit_resample(ind_train, dep_train)
```

```
[146]: ind_train_sampled.shape, dep_train_sampled.shape
```

```
[146]: ((15039, 8), (15039,))
```

```
[147]: print(f"The sampled training dataset is: {dep_train_sampled.value_counts()}")
print(f"The original training dataset is: {dep_train.value_counts()}")
```

The sampled training dataset is: 11      557

16      557

26      557

29      557

1      557

27      557

25      557

22      557

18      557

3      557

5      557

14      557

6      557

21      557

4      557

17      557

20      557

19      557

12      557

9      557

15      557

10      557

7      557

23      557

13      557

8      557

24      557

Name: Rings, dtype: int64

The original training dataset is: 9      557

10      491

8      461

11      396

7      311

6      206

12      204

13      164

14      102

5      95

```

15      82
16      57
17      46
4       46
18      32
19      25
20      22
3       13
21      11
23       7
22       5
27       2
24       2
25       1
1        1
29       1
26       1
Name: Rings, dtype: int64

```

## 2 Apply Z-Score normalization on the sampled dataset

```
[148]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
[149]: ind_train_sampled.iloc[:, 1:] = sc.fit_transform(ind_train_sampled.iloc[:, 1:])
```

```
[153]: ind_train_sampled, dep_train_sampled
```

```
[153]: (
    Sex    Length  Diameter    Height  Whole_weight  Sucked_weight \
0      M -0.403965 -0.443999 -0.489373    -0.796526    -0.791146
1      I -1.578054 -1.544685 -1.612751    -1.459856    -1.436658
2      I -0.022386 -0.017927 -0.662201    -0.740038    -0.643479
3      I  0.564658  0.408146  0.115523    -0.017800     0.329009
4      I  0.535306  0.550170  0.115523     0.018514     0.191890
...    ..      ...      ...      ...      ...      ...
15034  F  1.034294  1.224784  0.634005     1.340332     1.419630
15035  F  1.034294  1.224784  0.634005     1.340332     1.419630
15036  F  1.034294  1.224784  0.634005     1.340332     1.419630
15037  F  1.034294  1.224784  0.634005     1.340332     1.419630
15038  F  1.034294  1.224784  0.634005     1.340332     1.419630

    Viscera_weight  Shell_weight
0          -0.417408    -0.734652
1          -1.440575    -1.400995
2          -0.840371    -0.619599
3          -0.199884    -0.140215
4          -0.006530     0.152209

```

```

...
15034      1.012608      0.775408
15035      1.012608      0.775408
15036      1.012608      0.775408
15037      1.012608      0.775408
15038      1.012608      0.775408

```

```
[15039 rows x 8 columns],
```

```

0         11
1          4
2          8
3         11
4         13

```

```

..
15034     29
15035     29
15036     29
15037     29
15038     29

```

```
Name: Rings, Length: 15039, dtype: int64)
```

```
[ ]:
```

```
[ ]:
```