

657a-asg1-ques2

February 9, 2023

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import warnings

warnings.filterwarnings('ignore')
```

```
[3]: df = pd.read_csv("D:/UWaterloo/Data Knowledge and Modelling/Assignment 1/
↳ abalone.csv", names = ['Sex', 'Length', 'Diameter', 'Height', 'Whole_weight',
↳ 'Sucked_weight', 'Viscera_weight', 'Shell_weight',
↳ 'Rings'], sep = ',')
```

```
[4]: df
```


```
[4]:
```

	Sex	Length	Diameter	Height	Whole_weight	Sucked_weight	\
0	M	0.455	0.365	0.095	0.5140	0.2245	
1	M	0.350	0.265	0.090	0.2255	0.0995	
2	F	0.530	0.420	0.135	0.6770	0.2565	
3	M	0.440	0.365	0.125	0.5160	0.2155	
4	I	0.330	0.255	0.080	0.2050	0.0895	
...	
4172	F	0.565	0.450	0.165	0.8870	0.3700	
4173	M	0.590	0.440	0.135	0.9660	0.4390	
4174	M	0.600	0.475	0.205	1.1760	0.5255	
4175	F	0.625	0.485	0.150	1.0945	0.5310	
4176	M	0.710	0.555	0.195	1.9485	0.9455	
		Viscera_weight	Shell_weight	Rings			
0		0.1010	0.1500	15			
1		0.0485	0.0700	7			
2		0.1415	0.2100	9			
3		0.1140	0.1550	10			
4		0.0395	0.0550	7			
...				
4172		0.2390	0.2490	11			
4173		0.2145	0.2605	10			
4174		0.2875	0.3080	9			
4175		0.2610	0.2960	10			
4176		0.3765	0.4950	12			

[4177 rows x 9 columns]

```
[5]: # Check any relationship between Sex and Rings
plt.scatter(y=df['Sex'], x=df['Rings'])
plt.xlabel('Rings')
plt.ylabel('Sex')
plt.show()

# This feature is not a very helpful tool to predict the Rings and hence we can
↳ drop it.
```



```
[6]: X = df.iloc[:, 1:8] # Removing sex feature
y = df.iloc[:, 8]
list_scores = []

# Split the dataset into 80:20
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=1)
```

```

# Apply KNN classification
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)
knn_initial_score = knn.score(X_test, y_test)
knn_initial_score

```

[6]: 0.20933014354066987

```

[7]: # Balance the training dataset
from imblearn.over_sampling import RandomOverSampler
os = RandomOverSampler(random_state=1)
X_train_sampled, y_train_sampled = os.fit_resample(X_train, y_train)

```

```

[8]: # Apply Z-score
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train_sampled.iloc[:, 1:] = sc.fit_transform(X_train_sampled.iloc[:, 1:])

```

```

[9]: from sklearn.metrics import accuracy_score
from sklearn.model_selection import StratifiedKFold

def model_training(K):
    model_score = []
    accuracy_list = []
    # Split the 80% training dataset using KFold
    kf = StratifiedKFold(n_splits=5)

    for train_index, test_index in kf.split(X_train_sampled, y_train_sampled):
        X_train, X_test = X_train_sampled.iloc[train_index,:], X_train_sampled.
        ↪iloc[test_index,:]
        y_train, y_test = y_train_sampled[train_index] , ↪
        ↪y_train_sampled[test_index]

        # create model for every fold
        knn = KNeighborsClassifier(n_neighbors=K)
        knn.fit(X_train, y_train)
        model_score.append(knn.score(X_test, y_test))
        pred_values = knn.predict(X_test)


        acc = accuracy_score(pred_values , y_test)
        accuracy_list.append(acc)

    avg_accuracy = sum(accuracy_list)/K
    return avg_accuracy

```

```
[10]: for i in range(5, 15):  
      avg_acc = model_training(i)  
      print(i, avg_acc)  
      plt.scatter(x=i, y=avg_acc*100)  
      plt.xlabel("Values of K across all folds")  
      plt.ylabel("Mean Validation Accuracy")
```

```
5 0.8094967681792131  
6 0.6637197713314323  
7 0.5589758964139399  
8 0.4843660310118588  
9 0.4236027701517578  
10 0.3772862102788529  
11 0.338242179434318  
12 0.30673061430833964  
13 0.279478603049513  
14 0.2561674734636406
```



As per the graph and on the basis of the type of dataset, we consider $K = 10$ to be the best choice.

```
[11]: # Running KNN using K=10
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=1)
knn = KNeighborsClassifier(n_neighbors=10)
knn.fit(X_train, y_train)
knn_initial_score = knn.score(X_test, y_test)
knn_initial_score
```

[11]: 0.22607655502392343

```
[15]: #KNN improvement using Weighted KNN
import numpy as np
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=1)
knn = KNeighborsClassifier(metric='wminkowski', p=1, n_neighbors=10,
    weights='distance', metric_params={'w': np.random.random(X_train.shape[1])})
knn.fit(X_train, y_train)
knn_initial_score = knn.score(X_test, y_test)
print(knn_initial_score)
```

0.23684210526315788